

```

1 /Users/xingkaiwu/PycharmProjects/Zillow/venv/bin/python /
  Users/xingkaiwu/PycharmProjects/Zillow/model_lgb.py
2 Number of features for LightGBM: 69
3   cooling_id  bathroom_cnt  bedroom_cnt  quality_id
   floor1_sqft \
4 0           0           2.0           3.0           4.0
   NaN
5 1          -1           3.5           4.0           NaN
   NaN
6 2           0           3.0           2.0           4.0
   NaN
7 3           0           2.0           2.0           4.0
   NaN
8 4          -1           2.5           4.0           NaN
   NaN
9
10   finished_area_sqft_calc  floor1_sqft_unk
   base_total_area  fireplace_cnt \
11 0                    1684.0           NaN
   NaN           NaN
12 1                    2263.0           NaN
   NaN           NaN
13 2                    2217.0           NaN
   NaN           NaN
14 3                     839.0           NaN
   NaN           NaN
15 4                    2283.0           NaN
   NaN           NaN
16
17   bathroom_full_cnt  garage_cnt  garage_sqft  spa_flag
   heating_id \
18 0           2.0           NaN           NaN           NaN
   1
19 1           3.0           2.0           468.0           NaN
   -1
20 2           3.0           NaN           NaN           NaN
   1
21 3           2.0           NaN           NaN           NaN
   1
22 4           2.0           2.0           598.0           NaN
   -1
23
24   latitude  longitude  lot_sqft  pool_cnt
   pool_total_size  pool_unk_2 \
25 0  34280992.0 -118488536.0   7528.0           NaN
   NaN           NaN
26 1  33668120.0 -117677552.0   3643.0           NaN
   NaN           NaN
27 2  34136312.0 -118175032.0  11423.0           NaN
   NaN           NaN

```

28	3	33755800.0	-118309000.0	70859.0	NaN
		NaN	NaN		
29	4	33485644.0	-117700232.0	6000.0	1.0
		NaN	NaN		
30					
31		pool_unk_3	landuse_type_id	census_1	city_id
		neighborhood_id \			
32	0	NaN	230	60371068.0	12447.0
		31817.0			
33	1	NaN	230	60590524.0	32380.0
		NaN			
34	2	NaN	230	60374640.0	47019.0
		275411.0			
35	3	NaN	235	60372964.0	12447.0
		54300.0			
36	4	1.0	230	60590424.0	17686.0
		NaN			
37					
38		region_zip	room_cnt	bathroom_small_cnt	unit_cnt
		patio_sqft	year_built \		
39	0	96370.0	0.0	NaN	1.0
		NaN	1959.0		
40	1	96962.0	0.0	1.0	NaN
		NaN	2014.0		
41	2	96293.0	0.0	NaN	1.0
		NaN	1940.0		
42	3	96222.0	0.0	NaN	1.0
		NaN	1987.0		
43	4	96961.0	8.0	1.0	NaN
		NaN	1981.0		
44					
45		story_cnt	tax_structure	tax_parcel	tax_year
		tax_property \			
46	0	NaN	122754.0	360170.0	2015.0
		6735.879883			237416.0
47	1	NaN	346458.0	585529.0	2015.0
		10153.019531			239071.0
48	2	NaN	61994.0	119906.0	2015.0
		11484.480469			57912.0
49	3	NaN	171518.0	244880.0	2015.0
		3048.739990			73362.0
50	4	2.0	169574.0	434551.0	2015.0
		5488.959961			264977.0
51					
52		tax_overdue_flag	tax_overdue_year		census_2
		avg_garage_size \			
53	0	NaN	NaN	6.037107e+13	
		NaN			
54	1	NaN	NaN	NaN	NaN
		234.0			

55	2	NaN	NaN	6.037464e+13
		NaN		
56	3	NaN	NaN	6.037296e+13
		NaN		
57	4	NaN	NaN	6.059042e+13
		299.0		
58				
59		property_tax_per_sqft	location_1	location_2
		location_3	location_4	\
60	0	3.999929	-84207544.0	152769536.0 -
		24963276.0	93525260.0	
61	1	4.486531	-84009432.0	151345664.0 -
		25170656.0	92506896.0	
62	2	5.180190	-84038720.0	152311344.0 -
		24951204.0	93223828.0	
63	3	3.633778	-84553200.0	152064800.0 -
		25398700.0	92910300.0	
64	4	2.404275	-84214592.0	151185872.0 -
		25364472.0	92335760.0	
65				
66		missing_finished_area	missing_total_area	
		missing_bathroom_cnt_calc	\	
67	0	0.0		1.0
		0.0		
68	1	0.0		1.0
		0.0		
69	2	0.0		1.0
		0.0		
70	3	0.0		1.0
		0.0		
71	4	0.0		1.0
		0.0		
72				
73		derived_room_cnt	avg_area_per_room	
		derived_avg_area_per_room	\	
74	0	5.0	NaN	336
		.799988		
75	1	7.5	NaN	301
		.733337		
76	2	5.0	NaN	443
		.399994		
77	3	4.0	NaN	209
		.750000		
78	4	6.5	285.375	351
		.230774		
79				
80		region_zip-groupcnt	region_zip-lot_sqft-diff	
		region_zip-lot_sqft-percent	\	
81	0	14719.0	-13398.968750	
		-0.640273		

82	1	17682.0	-2715.032715
		-0.427024	
83	2	4422.0	-14927.021484
		-0.566490	
84	3	7293.0	-43346.804688
		-0.379550	
85	4	9875.0	-1155.377441
		-0.161470	
86			
87	region_zip-year_built-diff region_zip-		
	finished_area_sqft_calc-diff \		
88	0	-3.998413	
		-247.725464	
89	1	35.535156	526.538208
90	2	-12.917847	-173.867432
91	3	21.690186	-782.150757
92	4	0.695679	244.801147
93			
94	region_zip-finished_area_sqft_calc-percent region_zip-		
	tax_structure-diff \		
95	0	-0.128241	
		-50475.015625	
96	1	0.303225	
		213678.171875	
97	2	-0.072721	
		-236757.281250	
98	3	-0.482466	
		30903.765625	
99	4	0.120107	
		-50359.125000	
100			
101	region_zip-tax_structure-percent region_zip-tax_land-		
	diff \		
102	0	-0.291377	51026.
	421875		
103	1	1.609267	16302.
	671875		
104	2	-0.792490	-427605.
	093750		
105	3	0.219777	-129440.
	812500		
106	4	-0.228975	-195977.
	250000		
107			
108	region_zip-tax_land-percent region_zip-tax_property-		
	diff \		

```

109 0 0.273762 2047.
    035645
110 1 0.073182 6339.
    847656
111 2 -0.880721 1845.
    573242
112 3 -0.638259 -1337.
    844971
113 4 -0.425156 -2742.
    872070
114
115 region_zip-tax_property-percent region_zip-
    property_tax_per_sqft-diff \
116 0 0.436576
    1.521634
117 1 1.662618
    2.160548
118 2 0.191471
    1.178391
119 3 -0.304986
    0.830251
120 4 -0.333203
    -1.339566
121
122 region_zip-property_tax_per_sqft-percent year month
    quarter
123 0 0.613984 0 1
    1
124 1 0.928875 0 1
    1
125 2 0.294465 0 1
    1
126 3 0.296145 0 1
    1
127 4 -0.357805 0 1
    1
128 0 0.0276
129 1 -0.1684
130 2 -0.0040
131 3 0.0218
132 4 -0.0050
133 Name: logerror, dtype: float32
134 X_train shape: (131426, 69)
135 y_train shape: (131426,)
136 X_val shape: (33578, 69)
137 y_val shape: (33578,)
138 [0, 13, 21, 66, 67, 68]
139 /Users/xingkaiwu/PycharmProjects/Zillow/venv/lib/python3.7
    /site-packages/lightgbm/engine.py:148: UserWarning: Found
    `num_boost_round` in params. Will use it instead of

```

```

139 argument
140     warnings.warn("Found `{}` in params. Will use it instead
      of argument".format(alias))
141 /Users/xingkaiwu/PycharmProjects/Zillow/venv/lib/python3.7
      /site-packages/lightgbm/engine.py:153: UserWarning: Found
      `early_stopping_rounds` in params. Will use it instead of
      argument
142     warnings.warn("Found `{}` in params. Will use it instead
      of argument".format(alias))
143 /Users/xingkaiwu/PycharmProjects/Zillow/venv/lib/python3.7
      /site-packages/lightgbm/basic.py:1295: UserWarning:
      categorical_feature in Dataset is overridden.
144 New categorical_feature is [0, 13, 21, 66, 67, 68]
145     'New categorical_feature is {}'.format(sorted(list(
      categorical_feature))))
146 [LightGBM] [Warning] Met negative value in categorical
      features, will convert it to NaN
147 [LightGBM] [Warning] Met negative value in categorical
      features, will convert it to NaN
148 [LightGBM] [Warning] Met negative value in categorical
      features, will convert it to NaN
149 [LightGBM] [Info] Total Bins 10627
150 [LightGBM] [Info] Number of data: 131426, number of used
      features: 69
151 [LightGBM] [Info] Start training from score 0.007854
152 Train score: 5.033494924135775
153 Val score: 6.871374063401431
154 Train on all data + Make predictions
155 {'objective': 'regression', 'metric': 'mae', 'num_threads
      ': 4, 'boosting_type': 'gbdt', 'num_boost_round': 1250, '
      learning_rate': 0.003, 'num_leaves': 127, 'min_data': 150
      , 'min_hessian': 0.001, 'max_depth': -1, 'max_bin': 255, '
      sub_feature': 0.5, 'bagging_fraction': 0.7, 'bagging_freq
      ': 50, 'min_data_per_group': 100, 'cat_smooth': 15.0, '
      lambda_l1': 0.0, 'lambda_l2': 0.0, 'feature_fraction_seed
      ': 2, 'bagging_seed': 3}
156 lgb_X: (164299, 69)
157 lgb_y: (164299,)
158 [LightGBM] [Warning] Met negative value in categorical
      features, will convert it to NaN
159 [LightGBM] [Warning] Met negative value in categorical
      features, will convert it to NaN
160 [LightGBM] [Warning] Met negative value in categorical
      features, will convert it to NaN
161 [LightGBM] [Info] Total Bins 10687
162 [LightGBM] [Info] Number of data: 164299, number of used
      features: 69
163 [LightGBM] [Info] Start training from score 0.007873
164 score: 6.736742332537429
165 Start model 0 (2016)

```

```

166 Start model 0 (2017)
167 Length of submission DataFrame: 2985217
168 Submission header:
169   ParcelId  201610  201611  201612  201710  201711
      201712
170 0  10754147  0.0015  0.0015  0.0015  0.0063  0.0063  0.
      0063
171 1  10759547  0.0023  0.0023  0.0023  0.0042  0.0042  0.
      0042
172 2  10843547  0.0295  0.0295  0.0295  0.0393  0.0393  0.
      0393
173 3  10859147  0.0338  0.0338  0.0338  0.0309  0.0309  0.
      0309
174 4  10879947 -0.0002 -0.0002 -0.0002 -0.0004 -0.0004 -0.
      0004
175 /Users/xingkaiwu/PycharmProjects/Zillow/venv/lib/python3.7
    /site-packages/lightgbm/engine.py:148: UserWarning: Found
    `num_boost_round` in params. Will use it instead of
    argument
176   warnings.warn("Found `{}` in params. Will use it instead
    of argument".format(alias))
177 /Users/xingkaiwu/PycharmProjects/Zillow/venv/lib/python3.7
    /site-packages/lightgbm/basic.py:1295: UserWarning:
    categorical_feature in Dataset is overridden.
178 New categorical_feature is [0, 13, 21, 66, 67, 68]
179   'New categorical_feature is {}'.format(sorted(list(
    categorical_feature))))
180 Ensemble Training and Prediction
181 lgb_X: (164299, 69)
182 lgb_y: (164299,)
183 Start training model 0
184 [LightGBM] [Warning] Met negative value in categorical
    features, will convert it to NaN
185 [LightGBM] [Warning] Met negative value in categorical
    features, will convert it to NaN
186 [LightGBM] [Warning] Met negative value in categorical
    features, will convert it to NaN
187 [LightGBM] [Info] Total Bins 10687
188 [LightGBM] [Info] Number of data: 164299, number of used
    features: 69
189 [LightGBM] [Info] Start training from score 0.007873
190 Start training model 1
191 [LightGBM] [Info] Total Bins 10687
192 [LightGBM] [Info] Number of data: 164299, number of used
    features: 69
193 [LightGBM] [Info] Start training from score 0.007873
194 Start training model 2
195 [LightGBM] [Info] Total Bins 10687
196 [LightGBM] [Info] Number of data: 164299, number of used
    features: 69

```

```
197 [LightGBM] [Info] Start training from score 0.007873
198 Start training model 3
199 [LightGBM] [Info] Total Bins 10687
200 [LightGBM] [Info] Number of data: 164299, number of used
    features: 69
201 [LightGBM] [Info] Start training from score 0.007873
202 Start training model 4
203 [LightGBM] [Info] Total Bins 10687
204 [LightGBM] [Info] Number of data: 164299, number of used
    features: 69
205 [LightGBM] [Info] Start training from score 0.007873
206 model 0: 6.735792228190091
207 model 1: 6.736534614820393
208 model 2: 6.734958556252753
209 model 3: 6.735870739248984
210 model 4: 6.7349624343664525
211 Saved 5 LightGBM models to files.
212 Start model 0 (2016)
213 Start model 0 (2017)
214 Start model 1 (2016)
215 Start model 1 (2017)
216 Start model 2 (2016)
217 Start model 2 (2017)
218 Start model 3 (2016)
219 Start model 3 (2017)
220 Start model 4 (2016)
221 Start model 4 (2017)
222 Length of submission DataFrame: 2985217
223 Submission header:
224   ParcelId  201610  201611  201612  201710  201711
    201712
225 0  10754147 -0.0011 -0.0011 -0.0011  0.0022  0.0022  0.
    0022
226 1  10759547 -0.0020 -0.0020 -0.0020 -0.0008 -0.0008 -0.
    0008
227 2  10843547  0.0318  0.0318  0.0318  0.0414  0.0414  0.
    0414
228 3  10859147  0.0359  0.0359  0.0359  0.0347  0.0347  0.
    0347
229 4  10879947 -0.0017 -0.0017 -0.0017 -0.0017 -0.0017 -0.
    0017
230
231 Process finished with exit code 0
232
```



Feature Importance

